

# COOL C++ COMPILER TIDBITS

MATT GODBOLT @MATTGODBOLT



JUNE 5<sup>TH</sup> 2021

# HELLO!

# HELLO!

I'm accidentally a verb



# COMPILER EXPLORER

# COMPILERS ARE AMAZING!!

```
void LightObject(Renderable *object)
{
    auto world(object->GetMatrix());
    SetWorldMatrix(world);
    auto lMat = world * worldToLight;
    auto boundLs = TransformBounds(
        object->GetBounds(), lMat);
    if (!boundLs.Overlap(lightSphere))
        return;
    auto ms = pos * world.Invert();
    ...
}
```

```
void LightObject(Renderable *object)
{
    auto world(object->GetMatrix());
    SetWorldMatrix(world);
    auto lMat = world * worldToLight;
    auto boundLs = TransformBounds(
        object->GetBounds(), lMat);
    if (!boundLs.Overlap(lightSphere))
        return;
    auto ms = pos * world.Invert();
    ...
}
```



```
template<typename T>
auto thingy(T &&t)
{
    auto bongo = t.fronk();
    if (bongo < T::min())
    {
        return t.badger();
    } else {
        return t.smeagol();
    }
}
```

```
template<typename T>
auto thingy(T &&t)
{
    auto bongo = t.fronk();
    if (bongo < T::min())
    {
        return t.badger();
    } else {
        return t.smeagol();
    }
}
```

```
In instantiation of 'auto thingy(T&&)
[with T = Marmoset]':
error: 'struct Marmoset' has no member
        named 'fronk'
        auto bongo = t.fronk();
                           ^~~~~~
error: 'min' is not a member of 'Marmoset'
        if (bongo < T::min())
                           ~~~~~^~
error: 'struct Marmoset' has no member
        named 'badger'
        return t.badger();
                           ^~~~~~
...
...
```

```
int sumSquared(
    const vector<int> &v) {
    int res = 0;
    for (auto i : v)
        res += i * i;
    return res;
}
```

```
int sumSquared(
    const vector<int> &v) {
    int res = 0;
    for (auto i : v)
        res += i * i;
    return res;
}
```

.L4 :

```
    vmovdqu ymm2, YMMWORD PTR [rax]
    add rax, 32
    vpmulld ymm0, ymm2, ymm2
    vpaddsd ymm1, ymm1, ymm0
    cmp rax, rdx
    jne .L4
```

# ASSEMBLY!?

# ASSEMBLY?

ȝ̄f̄w̄z̄d̄r̄w̄h̄r̄p̄ȝ̄. ȝ̄f̄w̄z̄d̄r̄p̄ȝ̄  
ȝ̄f̄w̄z̄d̄h̄ūr̄ȝ̄. ȝ̄f̄w̄z̄d̄h̄ūr̄ȝ̄

*Ssolbergj, Wikimedia Commons CC SA 4.0*

# X86-64 ASSEMBLY

```
instr  
instr dest_operand  
instr dest_operand, source_operand  
instr dest_operand, source_operand, source_operand
```

# X86-64 ASSEMBLY

```
instr  
instr dest_operand  
instr dest_operand, source_operand  
instr dest_operand, source_operand, source_operand
```

```
ret ; return  
inc rax ; increment "rax"  
mov edx, 1234 ; set "edx" to the value 1234  
add rsi, rdi ; "rsi" += "rdi"  
vpaddd ymm1, ymm2, ymm0 ; "ymm1" = "ymm2" + "ymm0"
```

# INSTRUCTIONS

- mov
- movzx
- movsxd
- lea
- call
- ret
- jmp

- push
- pop
- cmp
- test
- je
- jne

- and
- xor
- add
- sub
- shl
- shr
- sar

# INSTRUCTIONS

- mov
- movzx
- movsxd
- lea
- call
- ret
- jmp

- push
- pop
- cmp
- test
- je
- jne

- and
- xor
- add
- sub
- shl
- shr
- sar

# INSTRUCTIONS

- mov
- movzx
- movsxd
- lea
- call
- ret
- jmp

- push
- pop
- cmp
- test
- je
- jne

- and
- xor
- add
- sub
- shl
- shr
- sar

# INSTRUCTIONS

- mov
- movzx
- movsxd
- lea
- call
- ret
- jmp

- push
- pop
- cmp
- test
- je
- jne

- and
- xor
- add
- sub
- shl
- shr
- sar

# INSTRUCTIONS

- mov
- movzx
- movsxd
- lea
- call
- ret
- jmp

- push
- pop
- **cmp**
- **test**
- **je**
- **jne**

- and
- xor
- add
- sub
- shl
- shr
- sar

# INSTRUCTIONS

- mov
- movzx
- movsxd
- lea
- call
- ret
- jmp

- push
- pop
- cmp
- test
- je
- jne

- and
- xor
- add
- sub
- shl
- shr
- sar

# INSTRUCTIONS

- mov
- movzx
- movsxd
- lea
- call
- ret
- jmp

- push
- pop
- cmp
- test
- je
- jne

- and
- xor
- add
- sub
- shl
- shr
- sar

And many, many more...

# OPERANDS

```
register ; e.g. rax, rbx, ecx...
constant ; e.g. 1234
<size> ptr [register] ; e.g. DWORD PTR [rax]
<size> ptr [register + offset] ; e.g. BYTE PTR [rcd + rsi]
<size> ptr [register + offset + register2 * (1,2,4,8)]
```

# OPERANDS

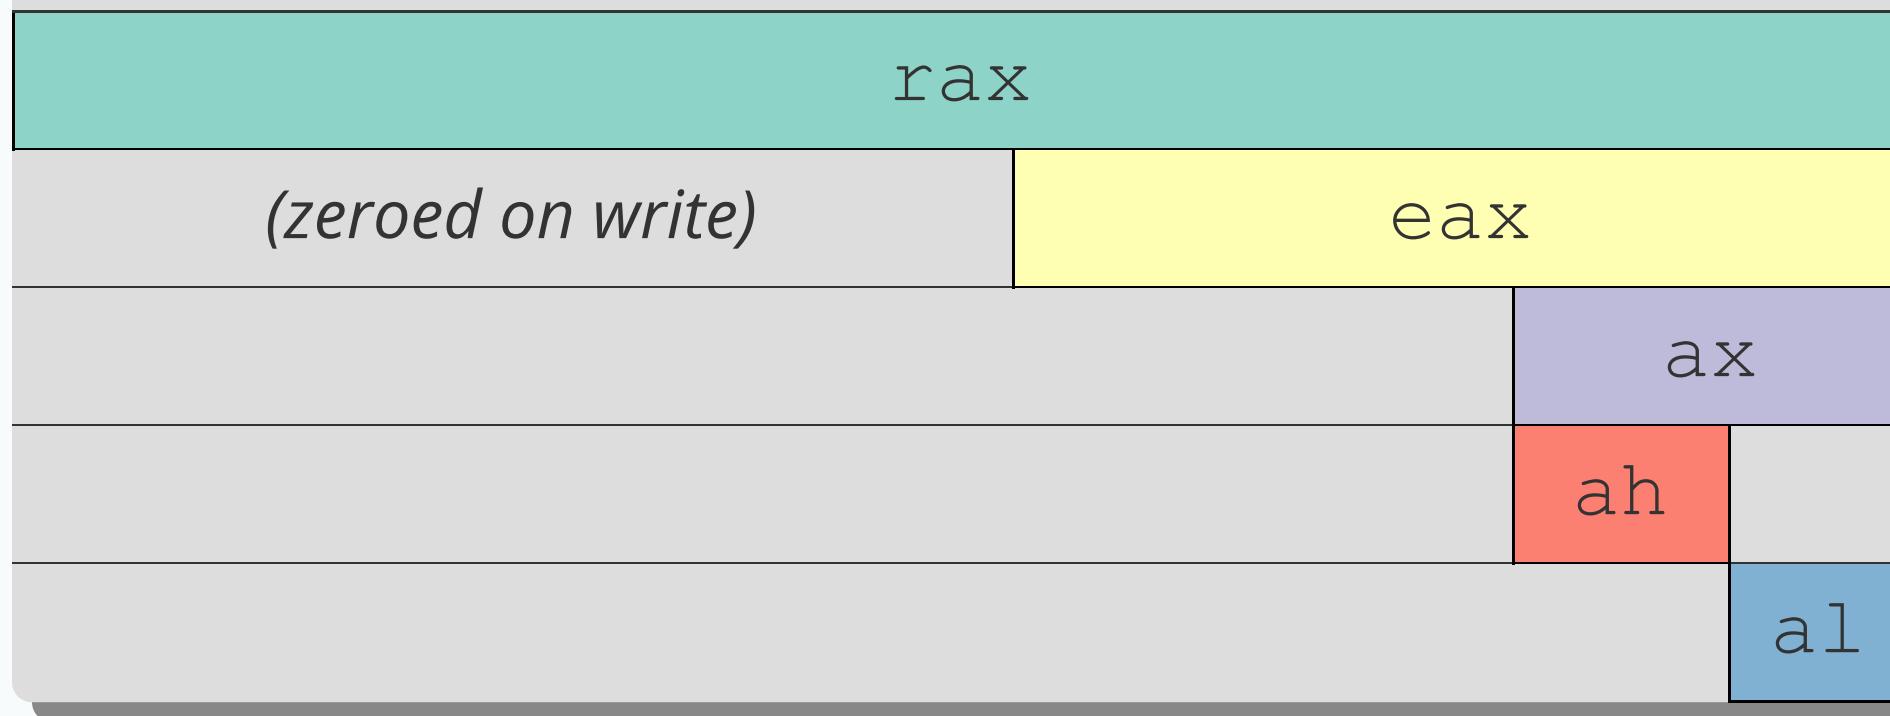
```
register ; e.g. rax, rbx, ecx...
constant ; e.g. 1234
<size> ptr [register] ; e.g. DWORD PTR [rax]
<size> ptr [register + offset] ; e.g. BYTE PTR [rcd + rsi]
<size> ptr [register + offset + register2 * (1,2,4,8)]
```

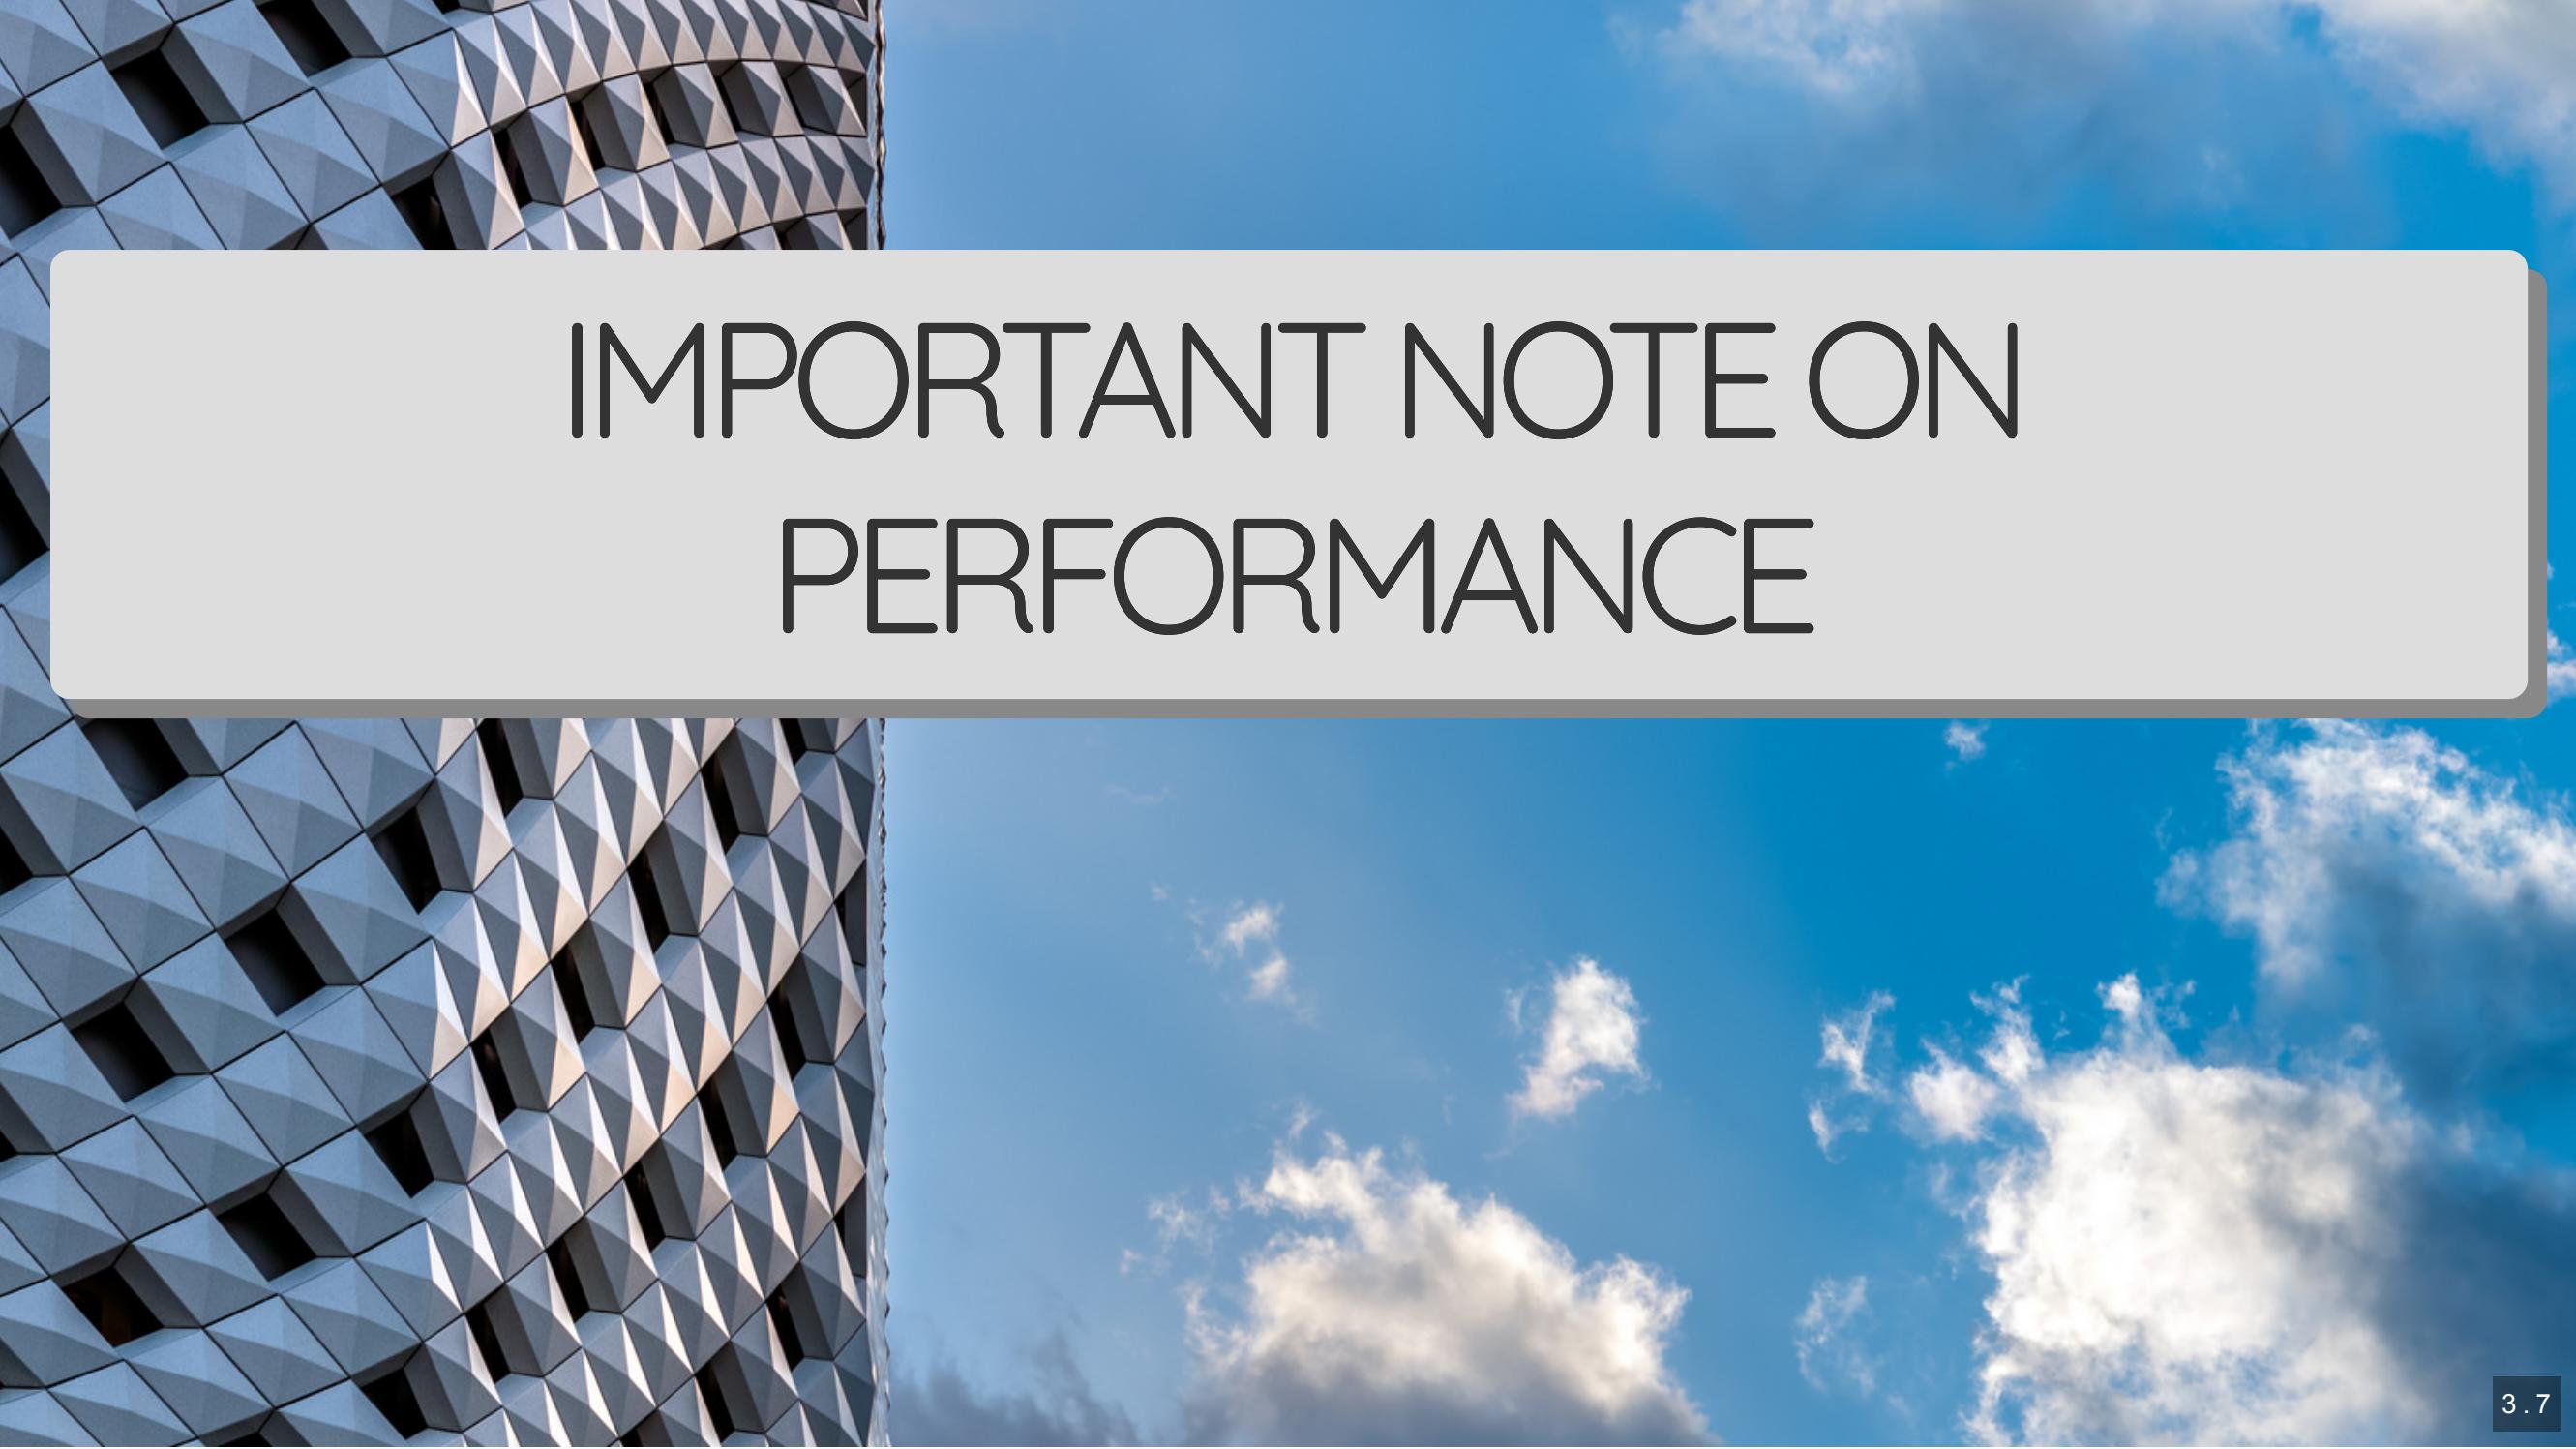
```
add eax, dword ptr [rdi + 12 + rsi * 4]
; eax += *(int *) (rdi + 12 + rsi * 4)
```

# REGISTERS

- `rax` (return value)
- `rdi` (1st param)
- `rsi` (2nd param)
- `rdx` (3rd param)
- `rbx` `rcx` `rbp` `r8-r15` `rsp`
- `xmm0-15` (`ymm0-15...` `zmm0-31...` `k0-7`)

**63...56    55...48    47...40    39...32    31...24    23...16    15...8    7...0**





# IMPORTANT NOTE ON PERFORMANCE

# IMPORTANT NOTE ON PERFORMANCE

- Google Benchmark
- hayai
- Nonius
- Quick Bench

# TIDBITS!

# MATH TIDBITS!

```
int test(int x, int y)
{
    return x - y;
}
```

```
int mulBy65599(int a)
{
    return (a << 16) + (a << 6) - a;
    //
    //      a * 65536      |
    //                      a * 64
    // 65536a + 64a - 1a = 65599a
}
```

```
int divideBy16(int a)
{
    return a >> 4; // clearly can't trust the compiler to do this for me
}
```

```
unsigned sumUpTo(unsigned x)
{
    auto total = 0u;
    for (auto i = 0u; i < x; ++i)
    {
        total += i;
    }
    return total;
}
```

```
unsigned sumUpTo(unsigned x)
{
    // equivalent to x(x+1)/2
    return ((x - 1) * (x - 2) / 2) + x - 1;
}
```

# VECTORIZATION TIDBITS!

```
int sumSquared(const vector<int> &v)
{
    int res = 0;
    for (auto i : v)
    {
        res += i * i;
    }
    return res;
}
```

```
int res_[] = {0,0,0,0,0,0,0,0};  
for (; index < v.size();  
      index += 8)  
{  
    for (size_t j = 0; j < 8; ++j)  
    { // but this happens in one instruction...  
        auto val = v[index + j];  
        res_[j] += val * val;  
    }  
}  
res = res_[0] + res_[1]  
    + res_[2] + res_[3]  
    + res_[4] + res_[5]  
    + res_[6] + res_[7];
```

```
template<typename T>
auto sumSquared(const T &v)
{
    return std::accumulate(
        begin(v), end(v),
        typename T::value_type(),
        [ ](auto x, auto y)
        {
            return x + y * y;
        } );
}
```

# ARCHITECTURAL TIDBITS!

```
int countSetBits(unsigned a)
{
    int count = 0;
    while (a != 0)
    {
        count++;
        // clear bottom set bit
        a &= (a - 1);
    }
    return count;
}
```

```
uint32_t switchBits(uint32_t x)
{
    auto first = x & 0xff;
    auto second = (x >> 8) & 0xff;
    auto third = (x >> 16) & 0xff;
    auto fourth = (x >> 24) & 0xff;
    return
        (first << 24)
        | (second << 16)
        | (third << 8)
        | fourth;
}
```

```
bool isspc(char c)
{
    return c == ' '
        || c == '\r'
        || c == '\n'
        || c == '\t';
}
```

```
movabs rax, 0x100002600 ; 0b10000000000000000000001001100000000
;           ^                                ^  ^
;           +- bit 32                  |  +- b10 & 9
;                           bit 13 -+
; shift right 'rdi' times
and eax, 1               ; pick the lowest set bit
```

# FUNCTION CALL TIDBITS!

```
unsigned identity(unsigned num); // just returns "num"

unsigned sumUpTo(unsigned x)
{
    auto total = 0u;
    for (auto i = 0u; i < x; ++i)
    {
        total += identity(i);
    }
    return total;
}
```

```
struct Func
{
    int operator() (int x)
    {
        return x;
    }
};

unsigned sumUpTo(unsigned x, Func &func)
{
    auto total = 0u;
    for (auto i = 0u; i < x; ++i)
        total += func(i);
    return total;
}
```

# CONCLUSION

# CONCLUSION

- Compilers are cleverer than we are

# CONCLUSION

- Compilers are cleverer than we are
  - if they have the right information

# CONCLUSION

- Compilers are cleverer than we are
  - if they have the right information
- Assembly isn't *that* scary

# CONCLUSION

- Trust your compiler
- Don't compromise readability

# CONCLUSION

- Trust your compiler
- Don't compromise readability
- Be aware of compiler limitations:
  - Visibility
  - (and aliasing)

# THANKS

Compiler Explorer is at <https://godbolt.org/>

Thanks for listening!

Thanks to the Compiler Explorer implementors.

All background images are by and © Romain Guy  
licensed [CC BY-NC-SA 2.0](#). Thanks Romain!

# THANKS

Compiler Explorer is at <https://godbolt.org/>

Thanks for listening!

Thanks to the Compiler Explorer implementors.

All background images are by and © Romain Guy  
licensed [CC BY-NC-SA 2.0](#). Thanks Romain!

**GO READ SOME ASSEMBLY!**